

# Assignment is below at the end

- <https://scikit-learn.org/stable/modules/tree.html>
- <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>
- [https://scikit-learn.org/stable/modules/generated/sklearn.tree.plot\\_tree.html](https://scikit-learn.org/stable/modules/generated/sklearn.tree.plot_tree.html)

```
In [144... import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = (20, 6)
plt.rcParams['font.size'] = 14
import pandas as pd
```

```
In [145... df = pd.read_csv('../data/adult.data', index_col=False)
```

```
In [146... golden = pd.read_csv('../data/adult.test', index_col=False)
```

```
In [147... golden.head()
```

Out[147]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	cap
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	

```
In [148... df.head()
```

Out[148]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	cap
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	

In [149... df.columns

Out[149]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss', 'hours-per-week', 'native-country', 'salary'], dtype='object')

In [150... from sklearn import preprocessing

In [151... enc = preprocessing.OrdinalEncoder()

In [152... transform\_columns = ['sex']  
non\_num\_columns = ['workclass', 'education', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'native-country']

In [153... pd.get\_dummies(df[transform\_columns]).head()

Out[153]:

	sex_Female	sex_Male
0	0	1
1	0	1
2	0	1
3	0	1
4	1	0

In [154... x = df.copy()  
  
x = pd.concat([x.drop(non\_num\_columns, axis=1),  
pd.get\_dummies(df[transform\_columns])], axis=1,  
  
x["salary"] = enc.fit\_transform(df[["salary"]])

In [155... x.head()

Out[155]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male
0	39	77516	13	2174	0	40	0.0	0	1
1	50	83311	13	0	0	13	0.0	0	1
2	38	215646	9	0	0	40	0.0	0	1
3	53	234721	7	0	0	40	0.0	0	1
4	28	338409	13	0	0	40	0.0	1	0

In [156... xt = golden.copy()  
  
xt = pd.concat([xt.drop(non\_num\_columns, axis=1),  
pd.get\_dummies(golden[transform\_columns])], axis=1,  
  
xt["salary"] = enc.fit\_transform(golden[["salary"]])

```
In [157... xt.salary.value_counts()
```

```
Out[157]: 0.0    12435
          1.0    3846
          Name: salary, dtype: int64
```

```
In [158... enc.categories_
```

```
Out[158]: [array(['<=50K.', '>50K.'], dtype=object)]
```

```
In [159... from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

## Choose the model of your preference: DecisionTree or RandomForest

```
In [160... model = RandomForestClassifier(criterion='entropy')
```

```
In [161... model = DecisionTreeClassifier(criterion='entropy', max_depth=None)
```

```
In [162... model.fit(x.drop(['fnlwgt', 'salary'], axis=1), x.salary)
```

```
Out[162]: DecisionTreeClassifier(criterion='entropy')
```

```
In [163... model.tree_.node_count
```

```
Out[163]: 8325
```

```
In [164... list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.feature_importances_))
```

```
Out[164]: [('age', 0.3225890802977448),
 ('education-num', 0.1607231138360648),
 ('capital-gain', 0.2268185889866726),
 ('capital-loss', 0.0789379219010595),
 ('hours-per-week', 0.15539844295130806),
 ('sex_Female', 0.05433887624806965),
 ('sex_Male', 0.0011939757790807106)]
```

```
In [165... list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.feature_importances_))
```

```
Out[165]: [('age', 0.3225890802977448),
 ('education-num', 0.1607231138360648),
 ('capital-gain', 0.2268185889866726),
 ('capital-loss', 0.0789379219010595),
 ('hours-per-week', 0.15539844295130806),
 ('sex_Female', 0.05433887624806965),
 ('sex_Male', 0.0011939757790807106)]
```

```
In [166... x.drop(['fnlwgt', 'salary'], axis=1).head()
```

```
Out[166]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	sex_Female	sex_Male
0	39	13	2174	0	40	0	1
1	50	13	0	0	13	0	1
2	38	9	0	0	40	0	1
3	53	7	0	0	40	0	1
4	28	13	0	0	40	1	0

```
In [167... set(x.columns) - set(xt.columns)
```

```
Out[167]: set()
```

```
In [168... list(x.drop('salary', axis=1).columns)
```

```
Out[168]: ['age',
            'fnlwgt',
            'education-num',
            'capital-gain',
            'capital-loss',
            'hours-per-week',
            'sex_Female',
            'sex_Male']
```

```
In [169... predictions = model.predict(xt.drop(['fnlwgt', 'salary'], axis=1))
predictionsx = model.predict(x.drop(['fnlwgt', 'salary'], axis=1))
```

```
In [170... from sklearn.metrics import (
            accuracy_score,
            classification_report,
            confusion_matrix, auc, roc_curve
        )
```

```
In [171... accuracy_score(xt.salary, predictions)
```

```
Out[171]: 0.8202813095018734
```

```
In [172... accuracy_score(xt.salary, predictions)
```

```
Out[172]: 0.8202813095018734
```

```
In [173... confusion_matrix(xt.salary, predictions)
```

```
Out[173]: array([[11460,   975],
                 [ 1951,  1895]])
```

```
In [174... print(classification_report(xt.salary, predictions))
```

	precision	recall	f1-score	support
0.0	0.85	0.92	0.89	12435
1.0	0.66	0.49	0.56	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

```
In [175... print(classification_report(xt.salary, predictions))
```

	precision	recall	f1-score	support
0.0	0.85	0.92	0.89	12435
1.0	0.66	0.49	0.56	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

```
In [176... accuracy_score(x.salary, predictionsx)
```

```
Out[176]: 0.8955806025613464
```

```
In [177... confusion_matrix(x.salary, predictionsx)
```

```
Out[177]: array([[24097,   623],
        [ 2777,  5064]])
```

```
In [178... print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841
accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

```
In [179... print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841
accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

For the following use the above `adult` dataset. Start with only numerical features/columns.

1. Show the RandomForest outperforms the DecisionTree for a fixed `max_depth` by training using the train set and `precision`, `recall`, `f1` on golden-test set.
2. For RandomForest or DecisionTree and using the `adult` dataset, systematically add new columns, one by one, that are non-numerical but converted using the feature-extraction techniques we learned. Show `[precision, recall, f1]` for each additional feature added.
3. Optional: Using `gridSearch` find the most optimal parameters for your model

Warning: this can be computationally intensive and may take some time.

- [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.GridSearchCV.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html)
- [https://scikit-learn.org/stable/modules/grid\\_search.html](https://scikit-learn.org/stable/modules/grid_search.html)

```
In [180]: adult = df.copy().drop(['fnlwgt'], axis=1)
adult.head()
```

Out[180]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	cap
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	

```
In [181]: gold = golden.copy().drop(['fnlwgt'], axis=1)
gold.head()
```

Out[181]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	cap
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	
4	18	?	Some-college	10	Never-married	?	Own-child	White	Female	0	

```
In [182]: adult.dtypes
```

Out[182]:

age	int64
workclass	object
education	object
education-num	int64
marital-status	object
occupation	object
relationship	object
race	object
sex	object

```
capital-gain      int64
capital-loss      int64
hours-per-week    int64
native-country    object
salary            object
dtype: object
```

## 1. Show the RandomForest outperforms the DecisionTree for a fixed max\_depth by training using the train set and precision, recall, f1 on golden-test set.

```
In [183... x1 = adult.copy()
x1 = pd.concat([x1.drop(non_num_columns, axis=1),
                pd.get_dummies(adult[transform_columns]), axis=1,])
x1["salary"] = enc.fit_transform(adult[["salary"]])
x1.head()
```

```
Out[183]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male
0	39	13	2174	0	40	0.0	0	1
1	50	13	0	0	13	0.0	0	1
2	38	9	0	0	40	0.0	0	1
3	53	7	0	0	40	0.0	0	1
4	28	13	0	0	40	0.0	1	0

```
In [184... t1 = gold.copy()
t1 = pd.concat([t1.drop(non_num_columns, axis=1),
                pd.get_dummies(gold[transform_columns]), axis=1,])
t1["salary"] = enc.fit_transform(gold[["salary"]])
t1.head()
```

```
Out[184]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male
0	25	7	0	0	40	0.0	0	1
1	38	9	0	0	50	0.0	0	1
2	28	12	0	0	40	1.0	0	1
3	44	10	7688	0	40	1.0	0	1
4	18	10	0	0	30	0.0	1	0

```
In [185... dt1 = DecisionTreeClassifier(criterion='entropy', max_depth=5)
rf1 = RandomForestClassifier(criterion='entropy', max_depth=5)
```

```
In [186... dt1.fit(x1.drop(['salary'],axis=1),x1.salary)
rf1.fit(x1.drop(['salary'],axis=1),x1.salary)
```

```
Out[186]: RandomForestClassifier(criterion='entropy', max_depth=5)
```

```
In [187... list(zip(x1.drop(['salary'], axis=1).columns, dt1.feature_importances_))
```

```
Out[187]: [('age', 0.27942432987238175),
('education-num', 0.21268599174904443),
('capital-gain', 0.3524074865400652),
('capital-loss', 0.04235939343563895),
('hours-per-week', 0.008875209903870702),
```

```
(('sex_ Female', 0.0),
 ('sex_ Male', 0.10424758849899901)])
```

```
In [188... list(zip(x1.drop(['salary'], axis=1).columns, rfl.feature_importances_))
```

```
Out[188]: [('age', 0.25060207498456627),
 ('education-num', 0.2100045738552291),
 ('capital-gain', 0.2834041907641028),
 ('capital-loss', 0.05455832340210596),
 ('hours-per-week', 0.0785636245268627),
 ('sex_ Female', 0.06803882879627765),
 ('sex_ Male', 0.05482838367085557)]
```

```
In [189... dtpred1 = dtl.predict(tl.drop(['salary'],axis=1))
 rfpred1 = rfl.predict(tl.drop(['salary'],axis=1))
```

```
In [190... accuracy_score(tl.salary, dtpred1)
```

```
Out[190]: 0.8200356243473989
```

```
In [191... accuracy_score(tl.salary, rfpred1)
```

```
Out[191]: 0.8341010994410663
```

```
In [192... confusion_matrix(tl.salary, dtpred1)
```

```
Out[192]: array([[11457,   978],
 [ 1952,  1894]])
```

```
In [193... confusion_matrix(tl.salary, rfpred1)
```

```
Out[193]: array([[12070,   365],
 [ 2336,  1510]])
```

```
In [194... print(classification_report(tl.salary, dtpred1))
```

	precision	recall	f1-score	support
0.0	0.85	0.92	0.89	12435
1.0	0.66	0.49	0.56	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

```
In [195... print(classification_report(tl.salary, rfpred1))
```

	precision	recall	f1-score	support
0.0	0.84	0.97	0.90	12435
1.0	0.81	0.39	0.53	3846
accuracy			0.83	16281
macro avg	0.82	0.68	0.71	16281
weighted avg	0.83	0.83	0.81	16281

**2. For RandomForest or DecisionTree and using the adult dataset, systematically add new columns, one by one, that are non-numerical but converted using the feature-extraction**



techniques we learned. Show [precision, recall, f1] for each additional feature added.

```
In [196... adult.head()
```

Out[196]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0

```
In [197... adult.workclass.value_counts()
```

Out[197]:

```
Private                22696
Self-emp-not-inc      2541
Local-gov             2093
?                     1836
State-gov             1298
Self-emp-inc          1116
Federal-gov           960
Without-pay           14
Never-worked           7
Name: workclass, dtype: int64
```

```
In [198... x2 = x1.copy()
x2['workclass'] = enc.fit_transform(adult[['workclass']])
x2.head()
```

Out[198]:

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass
0	39	13	2174	0	40	0.0	0	1	7.0
1	50	13	0	0	13	0.0	0	1	6.0
2	38	9	0	0	40	0.0	0	1	4.0
3	53	7	0	0	40	0.0	0	1	4.0
4	28	13	0	0	40	0.0	1	0	4.0

```
In [199... t2 = t1.copy()
t2['workclass'] = enc.fit_transform(gold[['workclass']])
t2.head()
```

Out[199]:

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass
0	25	7	0	0	40	0.0	0	1	4.0

1	38	9	0	0	50	0.0	0	1	4.0
2	28	12	0	0	40	1.0	0	1	2.0
3	44	10	7688	0	40	1.0	0	1	4.0
4	18	10	0	0	30	0.0	1	0	0.0

```
In [200... rf2 = rf1.fit(x2.drop(['salary'],axis=1),x2.salary)
rfpred2 = rf2.predict(t2.drop(['salary'],axis=1))
print(classification_report(t2.salary, rfpred2))
```

	precision	recall	f1-score	support
0.0	0.84	0.98	0.90	12435
1.0	0.82	0.38	0.52	3846
accuracy			0.83	16281
macro avg	0.83	0.68	0.71	16281
weighted avg	0.83	0.83	0.81	16281

```
In [201... adult.education.value_counts()
```

```
Out[201]: HS-grad      10501
Some-college  7291
Bachelors    5355
Masters      1723
Assoc-voc    1382
11th         1175
Assoc-acdm   1067
10th         933
7th-8th      646
Prof-school  576
9th          514
12th         433
Doctorate    413
5th-6th      333
1st-4th      168
Preschool    51
Name: education, dtype: int64
```

```
In [202... x3 = x2.copy()
x3['education'] = enc.fit_transform(adult[['education']])
x3.head()
```

```
Out[202]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education
0	39	13	2174	0	40	0.0	0	1	7.0	9.0
1	50	13	0	0	13	0.0	0	1	6.0	9.0
2	38	9	0	0	40	0.0	0	1	4.0	11.0
3	53	7	0	0	40	0.0	0	1	4.0	1.0
4	28	13	0	0	40	0.0	1	0	4.0	9.0

```
In [203... t3 = t2.copy()
t3['education'] = enc.fit_transform(gold[['education']])
t3.head()
```

```
Out[203]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education
0	25	7	0	0	40	0.0	0	1	4.0	1.0

1	38	9	0	0	50	0.0	0	1	4.0	11.0
2	28	12	0	0	40	1.0	0	1	2.0	7.0
3	44	10	7688	0	40	1.0	0	1	4.0	15.0
4	18	10	0	0	30	0.0	1	0	0.0	15.0

```
In [204... rf3 = rf1.fit(x3.drop(['salary'],axis=1),x3.salary)
rfpred3 = rf3.predict(t3.drop(['salary'],axis=1))
print(classification_report(t3.salary, rfpred3))
```

```

              precision    recall  f1-score   support

    0.0         0.84      0.97      0.90      12435
    1.0         0.78      0.39      0.52       3846

 accuracy              0.83      16281
 macro avg           0.81      0.68      0.71      16281
 weighted avg        0.82      0.83      0.81      16281
```

```
In [205... adult['marital-status'].value_counts()
```

```
Out[205]: Married-civ-spouse      14976
Never-married      10683
Divorced           4443
Separated          1025
Widowed            993
Married-spouse-absent  418
Married-AF-spouse    23
Name: marital-status, dtype: int64
```

```
In [206... x4 = x3.copy()
x4['marital-status'] = enc.fit_transform(adult[['marital-status']])
x4.head()
```

```
Out[206]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status
0	39	13	2174	0	40	0.0	0	1	7.0	9.0	4.0
1	50	13	0	0	13	0.0	0	1	6.0	9.0	2.0
2	38	9	0	0	40	0.0	0	1	4.0	11.0	0.0
3	53	7	0	0	40	0.0	0	1	4.0	1.0	2.0
4	28	13	0	0	40	0.0	1	0	4.0	9.0	2.0

```
In [207... t4 = t3.copy()
t4['marital-status'] = enc.fit_transform(gold[['marital-status']])
t4.head()
```

```
Out[207]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status
0	25	7	0	0	40	0.0	0	1	4.0	1.0	4.0
1	38	9	0	0	50	0.0	0	1	4.0	11.0	2.0
2	28	12	0	0	40	1.0	0	1	2.0	7.0	2.0
3	44	10	7688	0	40	1.0	0	1	4.0	15.0	2.0

4 18 10 0 0 30 0.0 1 0 0.0 15.0 4.0

```
In [208... rf4 = rf1.fit(x4.drop(['salary'],axis=1),x4.salary)
rfpred4 = rf4.predict(t4.drop(['salary'],axis=1))
print(classification_report(t4.salary, rfpred4))
```

	precision	recall	f1-score	support
0.0	0.86	0.96	0.91	12435
1.0	0.80	0.48	0.60	3846
accuracy			0.85	16281
macro avg	0.83	0.72	0.75	16281
weighted avg	0.84	0.85	0.83	16281

```
In [209... adult.occupation.value_counts()
```

Out[209]: Prof-specialty 4140  
Craft-repair 4099  
Exec-managerial 4066  
Adm-clerical 3770  
Sales 3650  
Other-service 3295  
Machine-op-inspct 2002  
? 1843  
Transport-moving 1597  
Handlers-cleaners 1370  
Farming-fishing 994  
Tech-support 928  
Protective-serv 649  
Priv-house-serv 149  
Armed-Forces 9  
Name: occupation, dtype: int64

```
In [210... x5 = x4.copy()
x5['occupation'] = enc.fit_transform(adult[['occupation']])
x5.head()
```

Out[210]:

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	39	13	2174	0	40	0.0	0	1	7.0	9.0	4.0	
1	50	13	0	0	13	0.0	0	1	6.0	9.0	2.0	
2	38	9	0	0	40	0.0	0	1	4.0	11.0	0.0	
3	53	7	0	0	40	0.0	0	1	4.0	1.0	2.0	
4	28	13	0	0	40	0.0	1	0	4.0	9.0	2.0	

```
In [211... t5 = t4.copy()
t5['occupation'] = enc.fit_transform(gold[['occupation']])
t5.head()
```

Out[211]:

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	25	7	0	0	40	0.0	0	1	4.0	1.0	4.0	
1	38	9	0	0	50	0.0	0	1	4.0	11.0	2.0	
2	28	12	0	0	40	1.0	0	1	2.0	7.0	2.0	

3	44	10	7688	0	40	1.0	0	1	4.0	15.0	2.0
4	18	10	0	0	30	0.0	1	0	0.0	15.0	4.0

```
In [212... rf5 = rf1.fit(x5.drop(['salary'],axis=1),x5.salary)
rfpred5 = rf5.predict(t5.drop(['salary'],axis=1))
print(classification_report(t5.salary, rfpred5))
```

	precision	recall	f1-score	support
0.0	0.86	0.96	0.91	12435
1.0	0.78	0.49	0.60	3846
accuracy			0.85	16281
macro avg	0.82	0.72	0.75	16281
weighted avg	0.84	0.85	0.83	16281

```
In [213... adult.relationship.value_counts()
```

```
Out[213]: Husband      13193
Not-in-family    8305
Own-child        5068
Unmarried        3446
Wife             1568
Other-relative    981
Name: relationship, dtype: int64
```

```
In [214... x6 = x5.copy()
x6['relationship'] = enc.fit_transform(adult[['relationship']])
x6.head()
```

```
Out[214]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	39	13	2174	0	40	0.0	0	1	7.0	9.0	4.0	
1	50	13	0	0	13	0.0	0	1	6.0	9.0	2.0	
2	38	9	0	0	40	0.0	0	1	4.0	11.0	0.0	
3	53	7	0	0	40	0.0	0	1	4.0	1.0	2.0	
4	28	13	0	0	40	0.0	1	0	4.0	9.0	2.0	

```
In [215... t6 = t5.copy()
t6['relationship'] = enc.fit_transform(gold[['relationship']])
t6.head()
```

```
Out[215]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	25	7	0	0	40	0.0	0	1	4.0	1.0	4.0	
1	38	9	0	0	50	0.0	0	1	4.0	11.0	2.0	
2	28	12	0	0	40	1.0	0	1	2.0	7.0	2.0	
3	44	10	7688	0	40	1.0	0	1	4.0	15.0	2.0	
4	18	10	0	0	30	0.0	1	0	0.0	15.0	4.0	

```
In [216... rf6 = rf1.fit(x6.drop(['salary'],axis=1),x6.salary)
rfpred6 = rf6.predict(t6.drop(['salary'],axis=1))
```

```
print(classification_report(t6.salary, rfpred6))
```

	precision	recall	f1-score	support
0.0	0.86	0.96	0.91	12435
1.0	0.80	0.49	0.61	3846
accuracy			0.85	16281
macro avg	0.83	0.73	0.76	16281
weighted avg	0.84	0.85	0.84	16281

```
In [217... adult.race.value_counts()
```

```
Out[217]: White          27816
Black          3124
Asian-Pac-Islander  1039
Amer-Indian-Eskimo  311
Other           271
Name: race, dtype: int64
```

```
In [218... x7 = x6.copy()
x7 = pd.concat([x7,
                pd.get_dummies(adult['race'])], axis=1)
x7.head()
```

```
Out[218]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	39	13	2174	0	40	0.0	0	1	7.0	9.0	4.0	
1	50	13	0	0	13	0.0	0	1	6.0	9.0	2.0	
2	38	9	0	0	40	0.0	0	1	4.0	11.0	0.0	
3	53	7	0	0	40	0.0	0	1	4.0	1.0	2.0	
4	28	13	0	0	40	0.0	1	0	4.0	9.0	2.0	

```
In [219... t7 = t6.copy()
t7 = pd.concat([t7,
                pd.get_dummies(gold['race'])], axis=1)
t7.head()
```

```
Out[219]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	25	7	0	0	40	0.0	0	1	4.0	1.0	4.0	
1	38	9	0	0	50	0.0	0	1	4.0	11.0	2.0	
2	28	12	0	0	40	1.0	0	1	2.0	7.0	2.0	
3	44	10	7688	0	40	1.0	0	1	4.0	15.0	2.0	
4	18	10	0	0	30	0.0	1	0	0.0	15.0	4.0	

```
In [220... rf7 = rf1.fit(x7.drop(['salary'],axis=1),x7.salary)
rfpred7 = rf7.predict(t7.drop(['salary'],axis=1))
print(classification_report(t7.salary, rfpred7))
```

	precision	recall	f1-score	support
0.0	0.86	0.96	0.91	12435
1.0	0.79	0.49	0.61	3846

accuracy			0.85	16281
macro avg	0.82	0.73	0.76	16281
weighted avg	0.84	0.85	0.84	16281

```
In [221]: list(zip(x7.drop(['salary'], axis=1).columns, rf7.feature_importances_))
```

```
Out[221]: [('age', 0.06950563975752137),
 ('education-num', 0.13984502540377203),
 ('capital-gain', 0.2025029163128069),
 ('capital-loss', 0.025746355853208815),
 ('hours-per-week', 0.04426595319204749),
 ('sex_Female', 0.020204875460454814),
 ('sex_Male', 0.010260486788901894),
 ('workclass', 0.001467851033114902),
 ('education', 0.026506516228131535),
 ('marital-status', 0.1920228653408138),
 ('occupation', 0.010984498542867831),
 ('relationship', 0.25581504899509583),
 (' Amer-Indian-Eskimo', 0.00016098402626058223),
 (' Asian-Pac-Islander', 0.0001028875473341934),
 (' Black', 0.0002810218158809956),
 (' Other', 9.996609648700666e-05),
 (' White', 0.0002271076053000775)]
```

```
In [222]: adult['native-country'].value_counts()
```

```
Out[222]: United-States      29170
 Mexico                643
 ?                     583
 Philippines           198
 Germany               137
 Canada               121
 Puerto-Rico           114
 El-Salvador           106
 India                 100
 Cuba                  95
 England               90
 Jamaica               81
 South                80
 China                 75
 Italy                 73
 Dominican-Republic   70
 Vietnam               67
 Guatemala             64
 Japan                 62
 Poland                60
 Columbia              59
 Taiwan                51
 Haiti                 44
 Iran                  43
 Portugal              37
 Nicaragua             34
 Peru                  31
 France                29
 Greece                29
 Ecuador               28
 Ireland               24
 Hong                  20
 Cambodia              19
 Trinidad&Tobago       19
 Laos                  18
 Thailand              18
 Yugoslavia            16
 Outlying-US (Guam-USVI-etc) 14
```

```
Honduras 13
Hungary 13
Scotland 12
Holand-Netherlands 1
Name: native-country, dtype: int64
```

```
In [223... x8 = x7.copy()
x8['native-country'] = enc.fit_transform(adult[['native-country']])
x8.head()
```

```
Out[223]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	39	13	2174	0	40	0.0	0	1	7.0	9.0	4.0	
1	50	13	0	0	13	0.0	0	1	6.0	9.0	2.0	
2	38	9	0	0	40	0.0	0	1	4.0	11.0	0.0	
3	53	7	0	0	40	0.0	0	1	4.0	1.0	2.0	
4	28	13	0	0	40	0.0	1	0	4.0	9.0	2.0	

```
In [224... t8 = t7.copy()
t8['native-country'] = enc.fit_transform(gold[['native-country']])
t8.head()
```

```
Out[224]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	salary	sex_Female	sex_Male	workclass	education	marital-status	occ
0	25	7	0	0	40	0.0	0	1	4.0	1.0	4.0	
1	38	9	0	0	50	0.0	0	1	4.0	11.0	2.0	
2	28	12	0	0	40	1.0	0	1	2.0	7.0	2.0	
3	44	10	7688	0	40	1.0	0	1	4.0	15.0	2.0	
4	18	10	0	0	30	0.0	1	0	0.0	15.0	4.0	

```
In [225... rf8 = rf1.fit(x8.drop(['salary'],axis=1),x8.salary)
rfpred8 = rf8.predict(t8.drop(['salary'],axis=1))
print(classification_report(t8.salary, rfpred8))
```

```

              precision    recall  f1-score   support

    0.0         0.86      0.96      0.91      12435
    1.0         0.79      0.48      0.60       3846

 accuracy          0.85      16281
 macro avg         0.82      0.72      0.75      16281
 weighted avg      0.84      0.85      0.83      16281
```

```
In [226... list(zip(x8.drop(['salary'], axis=1).columns, rf8.feature_importances_))
```

```
Out[226]: [('age', 0.09315313503748186),
('education-num', 0.141954314016355),
('capital-gain', 0.2013202078235829),
('capital-loss', 0.023589564343195653),
('hours-per-week', 0.04218309497075345),
('sex_Female', 0.024070322070793848),
('sex_Male', 0.023891420467357295),
('workclass', 0.002555344425558234),
('education', 0.025164904954484524),
```



```
(('marital-status', 0.17691583151348625),
 ('occupation', 0.010923233799191856),
 ('relationship', 0.23201689419975413),
 (' Amer-Indian-Eskimo', 6.575017624484862e-05),
 (' Asian-Pac-Islander', 7.217424952946028e-05),
 (' Black', 0.0007063204186603916),
 (' Other', 6.620254617335375e-05),
 (' White', 0.000435496119348729),
 ('native-country', 0.0009157888680481677)])
```

### 3. Optional: Using gridSearch find the most optimal parameters for your model

Warning: this can be computationally intensive and may take some time.

- [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.GridSearchCV.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html)
- [https://scikit-learn.org/stable/modules/grid\\_search.html](https://scikit-learn.org/stable/modules/grid_search.html)

```
In [227... rf.get_params().keys()
```

```
Out[227]: dict_keys(['bootstrap', 'ccp_alpha', 'class_weight', 'criterion', 'max_depth', 'max_features', 'max_leaf_nodes', 'max_samples', 'min_impurity_decrease', 'min_samples_leaf', 'min_samples_split', 'min_weight_fraction_leaf', 'n_estimators', 'n_jobs', 'oob_score', 'random_state', 'verbose', 'warm_start'])
```

```
In [228... from sklearn.model_selection import GridSearchCV
rf = RandomForestClassifier(criterion='entropy')
param_grid_rfc = [{
    'max_depth':[2, 3, 4, 5, 6, 7, 8],
    'max_features':[2, 3, 4, 5, 6, 7, 8]
}]
clf = GridSearchCV(estimator=rf,
                   param_grid = param_grid_rfc,
                   scoring='accuracy',
                   cv=10,
                   refit=True,
                   n_jobs=1)
clf.fit(x8.drop(['salary'],axis=1),x8.salary)
print(clf.best_score_)
print(clf.best_params_)
clfRFC = clf.best_estimator_
print('Test accuracy: %.3f' % clfRFC.score(t8.drop(['salary'],axis=1), t8.salary))

0.8560551116891307
{'max_depth': 8, 'max_features': 8}
Test accuracy: 0.857
```